

CS 211
Spring 2018
Midterm
3/06/2018

Name: Answer
NetID: _____

Time limit: 80 Minutes

Instructor: Abhishek Bhattacharjee

This exam contains 10 pages (including cover page) and 11 questions, for 100 points total.

Points Table (for grader use only)

Question	Points	Score
1	5	
2	4	
3	20	
4	10	
5	10	
6	15	
7	10	
8	6	
9	6	
10	6	
11	8	
Total:	100	

By signing below, I pledge that I have not violated the terms of Rutgers University's academic integrity policy¹. I have not consulted other students during the exam, I have not relied on notes or textbooks, I have not used a calculator, and I have not used any other manner of study and/or digital aids. Furthermore, I have not provided any aid to other students taking the exam during the examination period. I understand that violating the terms of this pledge will mean that I will be subject to disciplinary action at the university-level.

Student's signature: _____

¹<http://academicintegrity.rutgers.edu>

1. (5 points) First, consider the following declarations:

```
// declarations
int i,j,k[10]={2,4,6,8,10,12,14,16,18,20};
int *ip;
char a,b,c[10]={11,12,13,14,15,16,17,18,19,20};
char *cp;
void *vp;

typedef struct stst {
    int i;
    char *p;
} st;

// assignment statements
st x;
ip = &k[4];
vp = ip;
cp = &c[0];
x.p = cp+3;
x.i = 27;
```

Evaluate the following expressions (if the expression is an error, indicate that with "X"):

- (a) (1 point) $*(ip + 2) + c[3]$; 28
- (b) (1 point) $(*cp + 1)$; 12
- (c) (1 point) $*(ip) + 6$; 16
- (d) (1 point) $vp = \&x; ((st*)vp) - \> p$; X or 14
- (e) (1 point) $*(ip - 2)$; 6
2. (4 points) In the programming assignment 1, fourth problem(hash table), what was the strategy to deal with conflict resolution?

- A. Double hashing
- ☒ B. Linear probing
- C. Separate chaining with linked lists
- D. Cuckoo hashing


3. (20 points) For the following items, write **T** if the sentence is *True*, **F** if it is *False*, or **X** if you don't know the answer. You get +1 point for each correct answer, 0 points for an item answered with an X, -1 point for each wrong answer.

Assume x86 32-bit architecture.

- T The stack grows towards lower addresses.
- F Two's complement notation has two different representations of zero, this is the reason that its name is *two's complement*.
- F Accessing data from memory (*RAM*) is faster than from registers.
- T Dynamic memory regions, i.e., regions of the memory allocated by *malloc()*, are located in the heap.
- T The line of C code: *val = val << 4;* multiplies *val* by 16.
- T *0xDEADBEEF* = 1101 1110 1010 1101 1011 1110 1110 1111 b
- T The C *for* loop: *for(i = 0; i ≤ 10; i++)* iterates over the loop 11 times
- F In two's complement notation, the most significant bit of a negative number is zero.
- T IEEE floating-point representation uses signed magnitude format.
- F Arithmetic operations on floating point numbers are faster than on integers.
- T The mantissa field of a floating point number can be determined for normalized representation by assuming a single 1 at the left of the decimal point.
- F A Von Neumann machine stores data in memory and instructions on registers.
- T ISA stands for Instruction Set Architecture. It is the contract between software and hardware on what instructions are supported by the machine.
- F *c = (char **)malloc(sizeof(char*) * 4);* allocates 4 bytes of memory in the heap;
- T A pointer in C is a number that stores a memory address.
- F C is an interpreted language.
- F *%ebp* and *%esp* are registers that always point to the stack and *%ebp* ≤ *%esp*.
- F General purpose registers in a 32-bit machine are 32-bit in size, meaning that instructions' operands are also 32-bit in size. The maximum amount of memory that a 32-bit machine can address is, then, 32GB.
- T Two's complement -1 in binary is: 1111 1111 1111 1111 1111 1111 1111 1111 b
- F *char*, *int*, *long*, *float*, *double*, *void*, and *string* are data types in C.

4. (10 points) Consider the code below. You will see a struct, and also a function, *printList*, that prints the linked list from its head to its tail. Your task is to implement *printReverseList*, a function that prints the list from its tail to its head.

Example, for the following linked list:



printList prints: 12 → 99 → 37 → 211 → 1 → NULL

printReverseList should print: 1 → 211 → 37 → 99 → 12 → NULL

```
typedef struct node_ {
    int data;
    struct node_ *next;
} node;
void printList(node *head) {
    while(head) {
        printf("%d-->", head->data);
        head = head->next;
    }
    printf("NULL\n");
}
void printReverseList(node *head){
    if (head == NULL)
        return;
    printReverseList(head->next);
    printf("%d-->", head->data);
```

← Recursive

-
- Actually reversing the list is also okay!
 - Using array is also okay!
 - Other... if it works I gave you points
and also you've got some partial credit
-

5. (10 points) What is the output of the following program?

```
#include <stdio.h>
#include <stdlib.h>
#define SIZE (5)
void matTr(int **mat, int n) {
    int i, j;
    for(i = 0; i < n; i++) {
        for(j = i+1; j < n; j++) {
            int tmp = mat[i][j];
            mat[i][j] = mat[j][i];
            mat[j][i] = tmp;
        }
    }
}

int main(int argc, char** argv) {
    int **mat, i, j;
    mat = (int **) malloc(sizeof(*mat) * SIZE);
    for(i = 0; i < SIZE; i++){
        mat[i] = (int *) malloc(sizeof(**mat) * SIZE);
        for(j = 0; j < SIZE; j++)
            mat[i][j] = (i*2 + j) % SIZE;
    }
    matTr(mat, SIZE);
    for(i = 0; i < n; i++) {
        for(j = 0; j < n; j++)
            printf("%d ", mat[i][j]);
        printf("\n");
    }
    for(i = 0; i < SIZE; i++)
        free(mat[i]);
    free(mat);
    return 0;
}
```

Output:

0	2	4	1	3
1	3	0	2	4
2	4	1	3	0
3	0	2	4	1
4	1	3	0	2

6. (15 points) For the following questions, assume the numbers are 16 bits.

1. Convert two's complement $0xC521$ to:

(a) (2 points) Decimal: -15071

(b) (2 points) Octal: 142441

(c) (2 points) Binary: $1100\ 0101\ 0010\ 0001$

2. Convert -512 (decimal) into hexadecimal with

(a) (3 points) Sign and magnitude representation: $0xB200$

(b) (3 points) One's complement: $0xFDFF$

(c) (3 points) Two's complement: $0xFE00$

7. (10 points) Implement a function that prints the hexadecimal representation of a decimal number (16-bit, two's complement). You are not allowed to use any library API or `printf()` specifier.

```
#include <stdio.h>
```

```
void dec2hex(short int num) {
    char hexa[5];
```

```
    int i, j = 0;
```

```
    while (num != 0) {
```

```
        int temp = 0;
```

```
        temp = num num % 16;
```

```
        if (temp > 9)
```

```
            hexa[i] = temp - 10 + 'A';
```

```
        else
```

```
            hexa[i] = temp + '0';
```

```
        num = num / 16;
```

```
        i++;
```

```
    }
```

```
    for (j = 0, --i; j < i; j++, i--) {
```

```
        char tmp = hexa[i];
```

```
        hexa[i] = hexa[j];
```

```
        hexa[j] = tmp;
```

```
    }
    printf("Hexa of %d is 0x%s.\n", num, hexa);
```

```
}
```

8. (6 points) As an application of the property that $a \oplus a = 0$ (\oplus is symbol for XOR) for any bit vector a , consider the following program:

```
void inplace_swap(int *x, int *y) {
    *y = *x ^ *y; // step 1
    *x = *x ^ *y; // step 2
    *y = *x ^ *y; // step 3
}
```

As the name implies, the effect of this procedure is to swap the values stored at the locations denoted by the pointer variables x and y . Note that unlike the usual technique for swapping two values, we do not need a third location to temporarily store one value while we are moving the other.

Starting with values a and b in the locations pointed by x and y , respectively, fill the table that follows, giving the values stored at the two locations after each step of the procedure. Use the properties of *XOR* to show that the desired effect is achieved. Recall that every element is its own additive inverse (that is, $a \oplus a = 0$).

step	*x	*y
Initially	a	b
step 1	a	$a \wedge b$
step 2	$(a \wedge a) \wedge b = b$	$a \wedge b$
step 3	b	$a \wedge b \wedge b = a$

9. (6 points) Consider the content of three files below.

File1	File2	File3
1 #include <stdio.h>	1 #ifdef __CS211__	1 C=gcc
2	2 #define __CS211__	2 CFLAGS=-I.
3 int main(void)	3	3
4 {	4 #define ABC (100)	4 hello: hello.o
5 printf("CS211\n");	5 typedef struct abc {	5 \$(CC) -o hello ↵
6 return 0;	6 int a, b, c;	6 hello.o -I.
7 }	7 };	
	8	
	9 #endif	

Identify which category each one of the files belong to:

(a) (2 points) Header file: File 2

(b) (2 points) Source file: File 1

(c) (1 point) Makefile: File 3

(d) (1 point) What is the simplest command-line statement to compile this code?

make

10. (6 points) Write below the function *CS211* that takes integer n as input and iterates over all the numbers from 1 up to n (inclusive). If the number is divisible by 3, then print "*Rutgers*". If the number is divisible by 5, then print "*CS211*". If the number is divisible by 3 and 5, then print "*Rutgers CS211*". Otherwise, print the number. For example, *CS211*(16) prints as follows:

```
1
2
Rutgers
4
CS211
Rutgers
7
8
Rutgers
CS211
11
Rutgers
13
14
Rutgers CS211
16
```

```
void CS211(int n) {
```

```
    for (int i = 1; i <= n; i++) {
        if (i % 3 == 0 && i % 5 == 0)
            printf("Rutgers CS 211\n");
        else if (i % 3 == 0)
            printf("Rutgers\n");
        else if (i % 5 == 0)
            printf("CS 211\n");
        else
            printf("%d\n");
    }
```

```
}
```

11. (8 points) Considering the following assembly code, answer the questions below.

```
1 .section .text
2 .globl main
3 main:    push %ebp          # stack frame setup
4          mov %esp, %ebp     # stack frame setup
✓ 5          mov $0x5, %ecx
✓ 6          mov %ecx, %edx
✓ 7 loop:   cmp $0x1, %ecx
✓ 8          jle end
✓ 9          dec %ecx
✓ 10         imul %ecx, %edx
✓ 11         jmp loop
✓ 12 end:    mov %edx, %eax
13          mov %ebp, %esp     # stack frame cleanup
14          pop %ebp          # stack frame cleanup
15          ret
```

We haven't covered in class yet how to setup and cleanup the stack frame (activation record) of a function, you won't need this information to answer this question. You can safely ignore the lines 1 – 4, and 13 – 15 to answer these questions.

Answer the questions:

- (a) (2 points) Content of `%eax` (in decimal) after executing the code: 120
- (b) (2 points) Number of times line 7 is executed: 5
- (c) (2 points) Number of memory references(ignore lines 1 – 4, 13 – 15): 25 or 30
- (d) (2 points) Concisely explain what this code does: _____

factorial(5), 5!, fact(5)